




DINOv3 encoder
action $R^2 = 0.30$



Ours (learned tokenizer)
action $R^2 = 0.21$



per-source standardised
(post hoc)



● robot demos ● human video