

DINOv3 encoder
action $R^2 = 0.30$


latent as trained

source kNN = 1.00




Ours (learned tokenizer)
action $R^2 = 0.21$

source kNN = 0.95




per-source standardised
(post hoc)

source kNN = 0.84



source kNN = 0.77



● robot demos

● human video