

DINOv3 encoder
action $R^2 = 0.30$


Ours (learned tokenizer)
action $R^2 = 0.21$

raw latent (8-D)

source kNN = 1.00




source kNN = 0.95




motion subspace (3-D)

source kNN = 0.76




source kNN = 0.70



pose directions removed (8-D)

source kNN = 0.83



source kNN = 0.64



● robot demos ● human video